

Project3

Buffer management

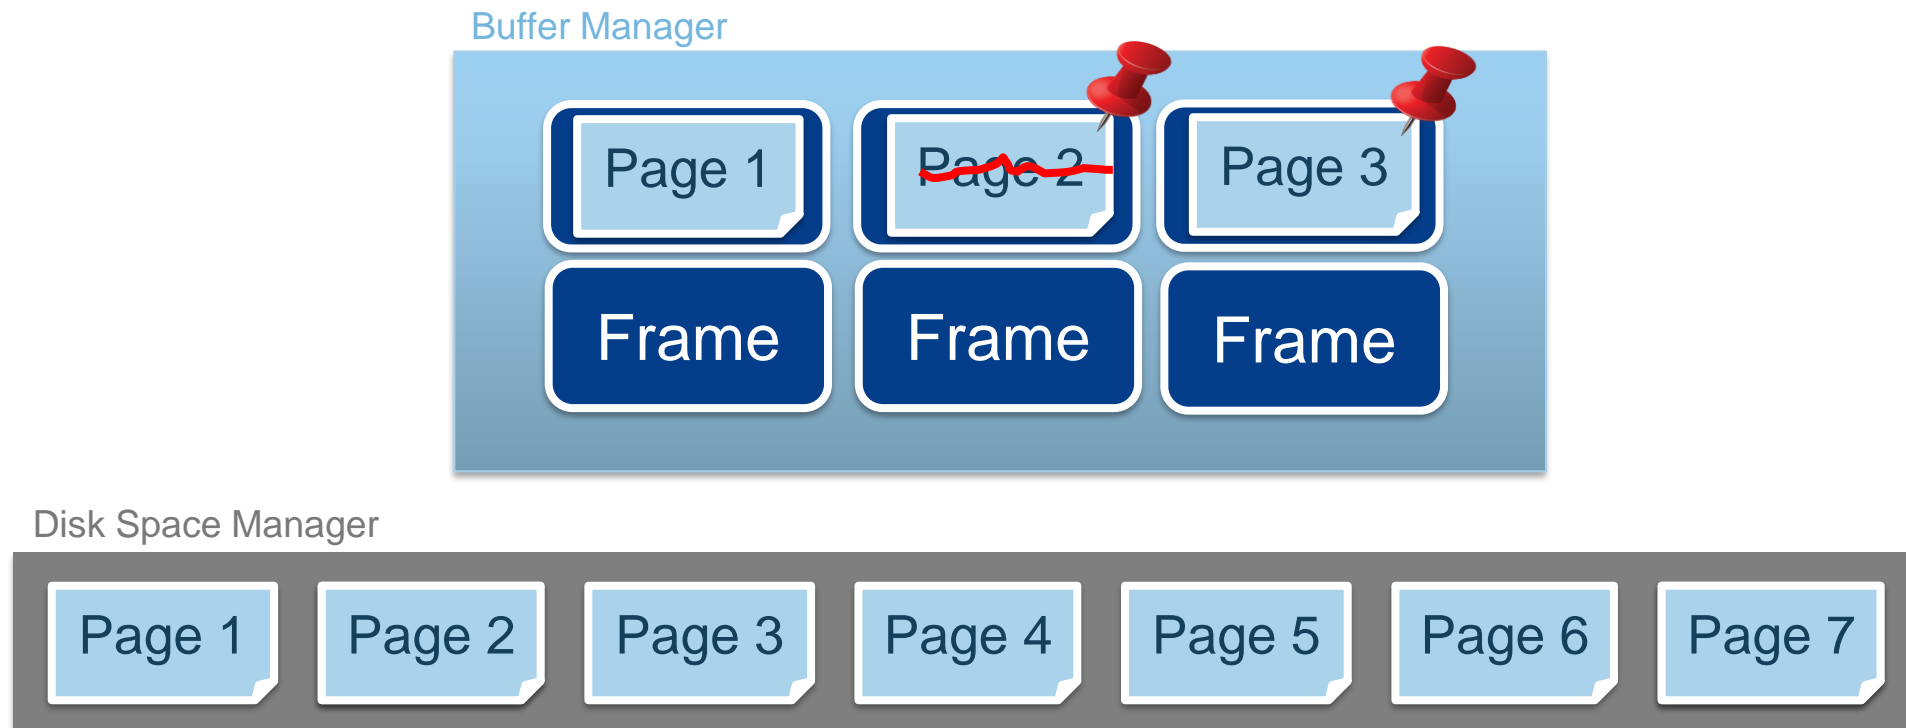
Submission

- You should submit your project in a directory structure like this: "your_repo/project3/db_project".
- Follow the db_project directory structure as before.

```
ktlee20@multicore-36:~/db2021$ ls
project1 project2 project3
ktlee20@multicore-36:~/db2021$ tree project3
project3
├── db_project
│   ├── CMakeLists.txt
│   ├── db
│   │   ├── CMakeLists.txt
│   │   ├── include
│   │   │   ├── bpt.h
│   │   │   ├── buffer.h
│   │   │   └── file.h
│   │   └── src
│   │       ├── bpt.cc
│   │       ├── buffer.c
│   │       └── file.cc
│   ├── DbConfig.h.in
│   ├── main.cc
│   └── test
│       ├── basic_test.cc
│       ├── CMakeLists.txt
│       └── file_test.cc
└── 5 directories, 13 files
```

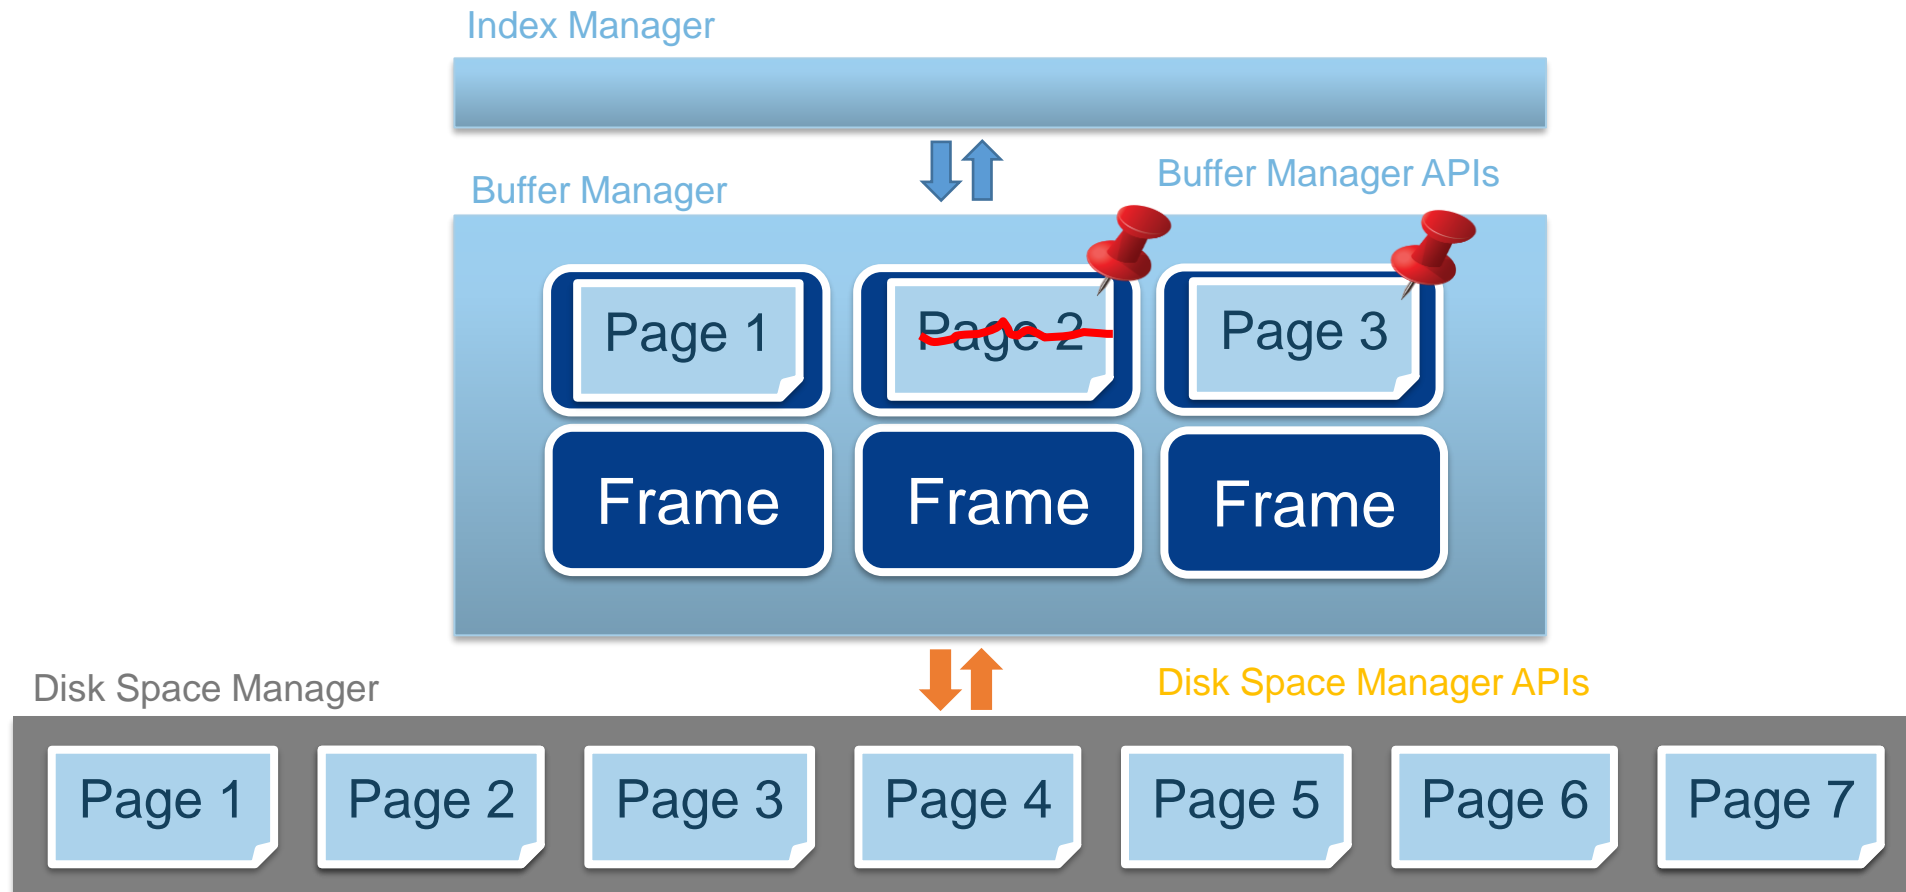
Buffer Management

- The current disk-based b+tree doesn't support buffer management.
- Our goal is to implement an **in-memory buffer manager** for caching on-disk pages.



Buffer Management

- Disk manager APIs should be called only within the buffer manager layer.



Project Specification

- Define a buffer block structure which contains the following fields.
- **Physical frame:** containing up-to-date contents of a target page.
 - **Table id:** the unique id of a table (per file)
 - **Page number:** the target page number within a file.
 - **Is dirty:** whether this buffer block is dirty or not.
 - **Is pinned:** whether this buffer is-use(pinned) or not.
 - **LRU list next/prev :** data for representing a LRU list
 - Other data necessary for your design.

Buffer Structure

frame (page size : 4096 bytes)
table_id
page_num
is_dirty
is_pinned
next of LRU
prev of LRU

Project Specification

- Modify the database initialization function.
 - **int init_db (int num_buf);**
 - Allocate the buffer pool with the given number of entries (i.e., num_buf).
 - Initialize other fields for your own design.
 - On success, return 0. Otherwise, return a non-zero value.

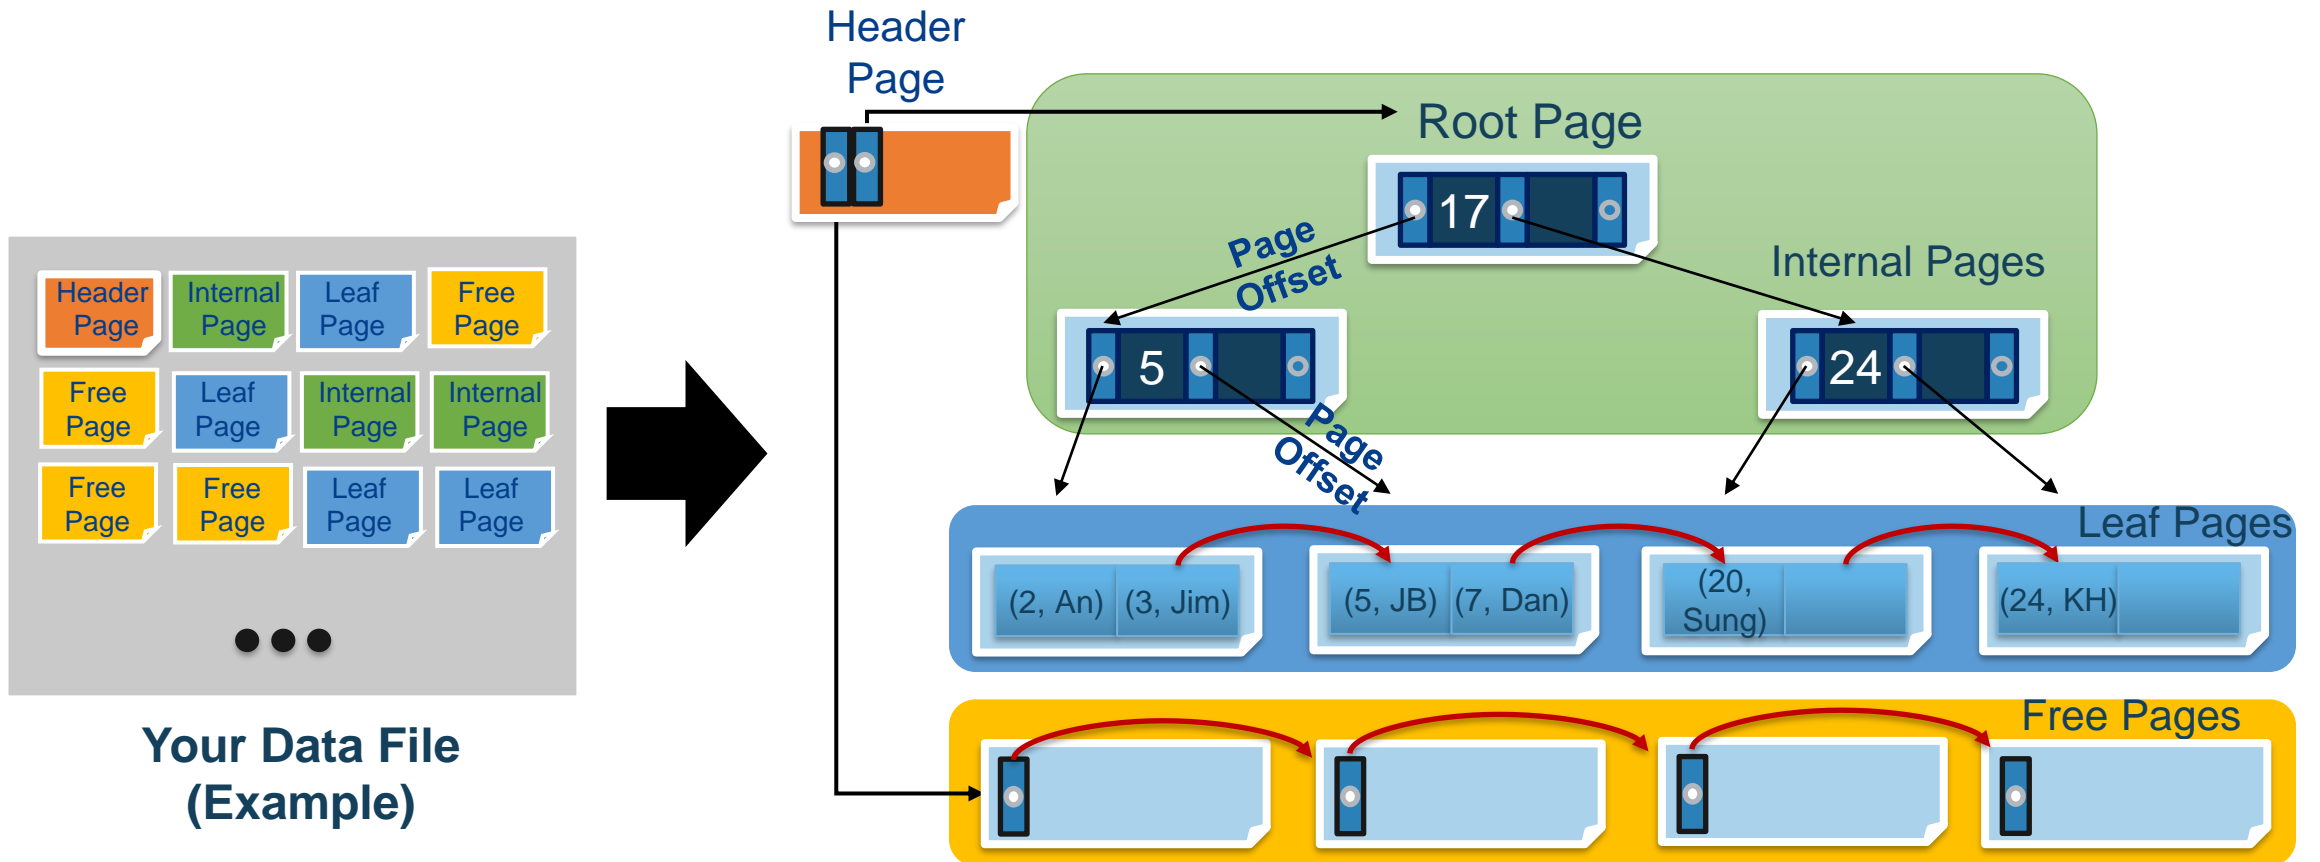
Project Specification

- Your APIs must always go through the buffer manager layer for accessing any data. (more details in the following slides)
 - If a page is not in the buffer pool (i.e., cache-miss), read the page from disk and maintain it in a buffer block.
 - Page modification should only occur within the in-memory buffer. If a page frame in memory is updated, mark the buffer block as dirty.
 - Select the victim page for eviction by following the LRU policy. Write the page to disk during the eviction process.

Project Specification

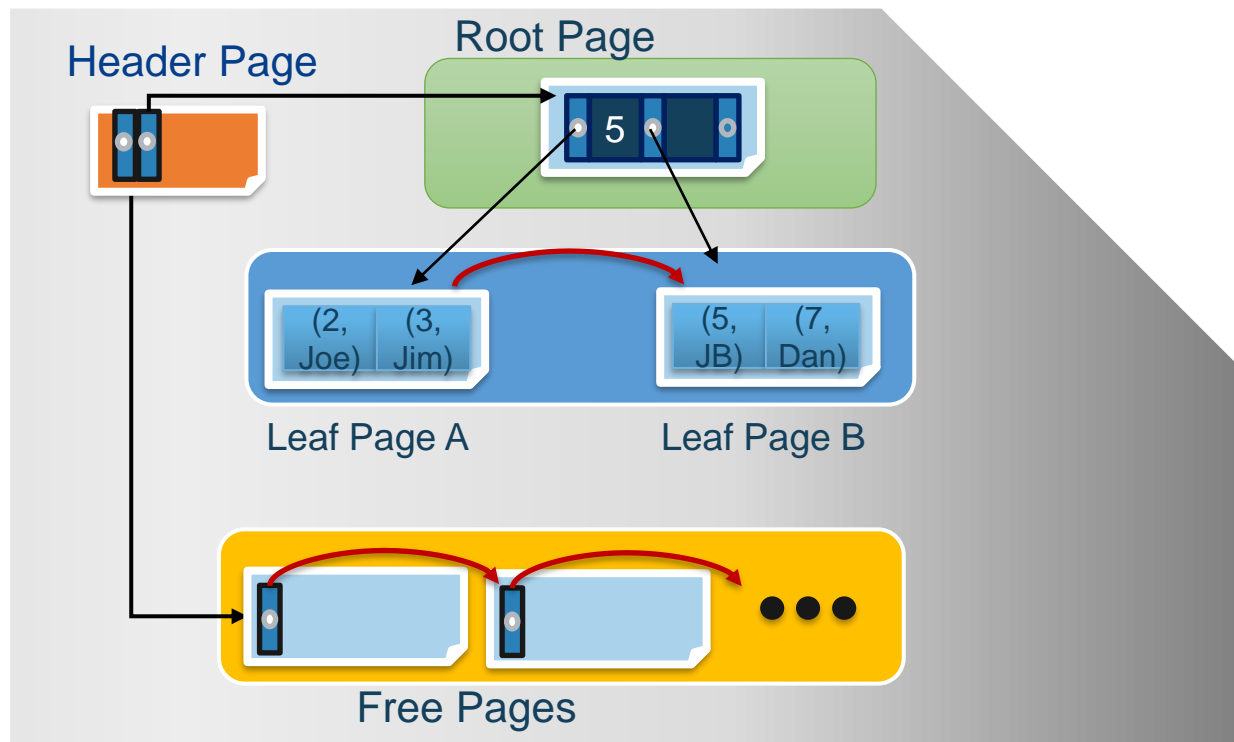
- Modify database shutdown function.
 - **int shutdown_db();**
 - Flush the entire buffer and destroy(e.g., deallocate, etc.) it.
 - On success, return 0. Otherwise, return a non-zero value.
 - Call file_close_table_files() to dereference the open files.

So far..




Buffer Management

- Assume that on-disk pages are stored like below.




Client

Buffer Manager

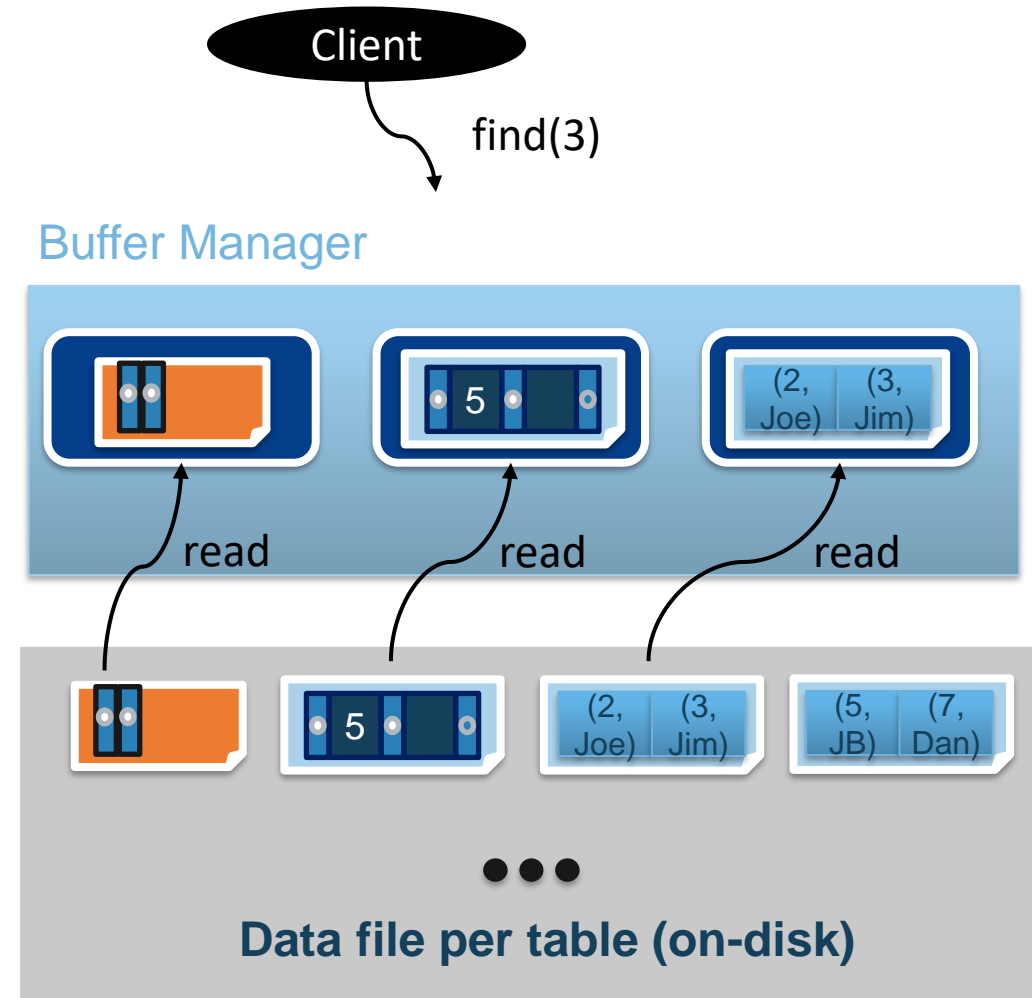


Disk Space Manager



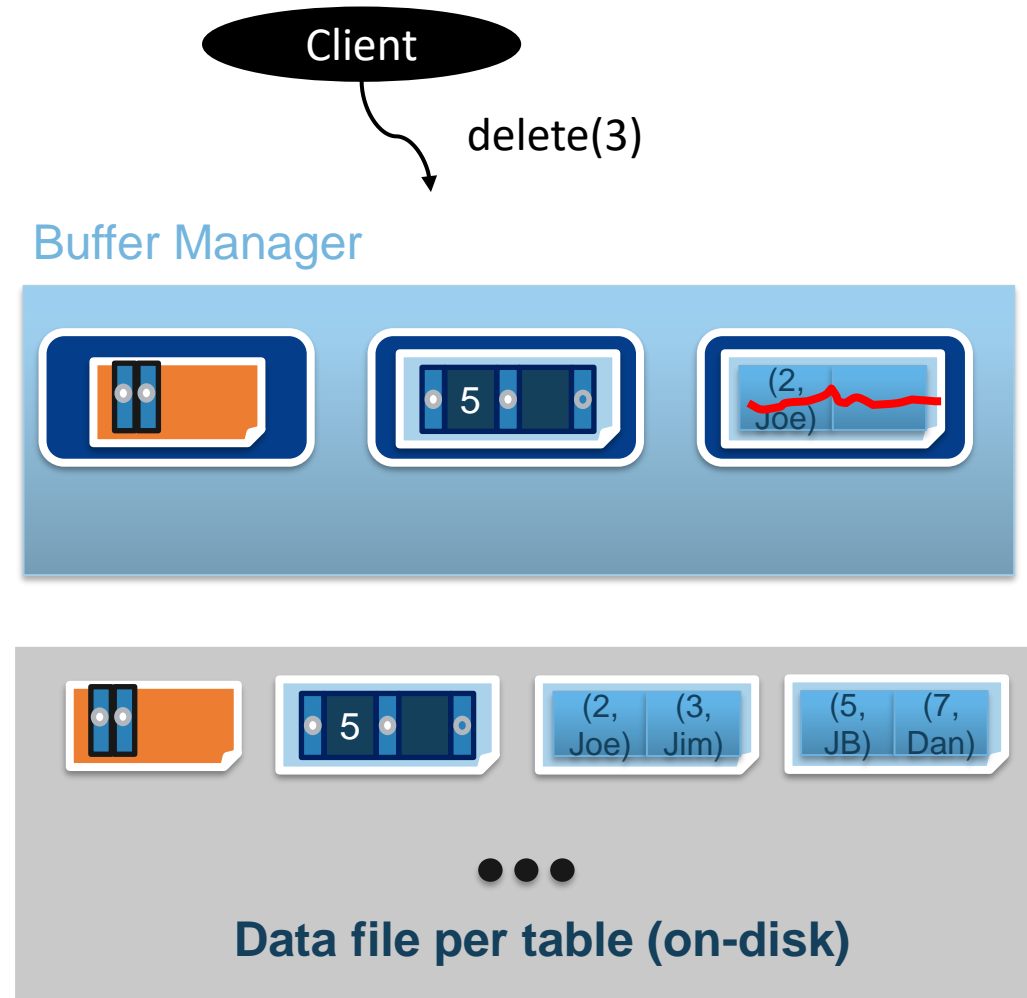
Buffer Management

- First, search for the target page within the buffer pool.
- If the desired page is not found (i.e., cache-miss occurs), read the page from disk and maintain it in a buffer block.
- Traversing the index from root to leaf page incurs access to various page types (e.g., header, root, internal, leaf, etc.). All pages are to be read/modified through the buffer manager layer.



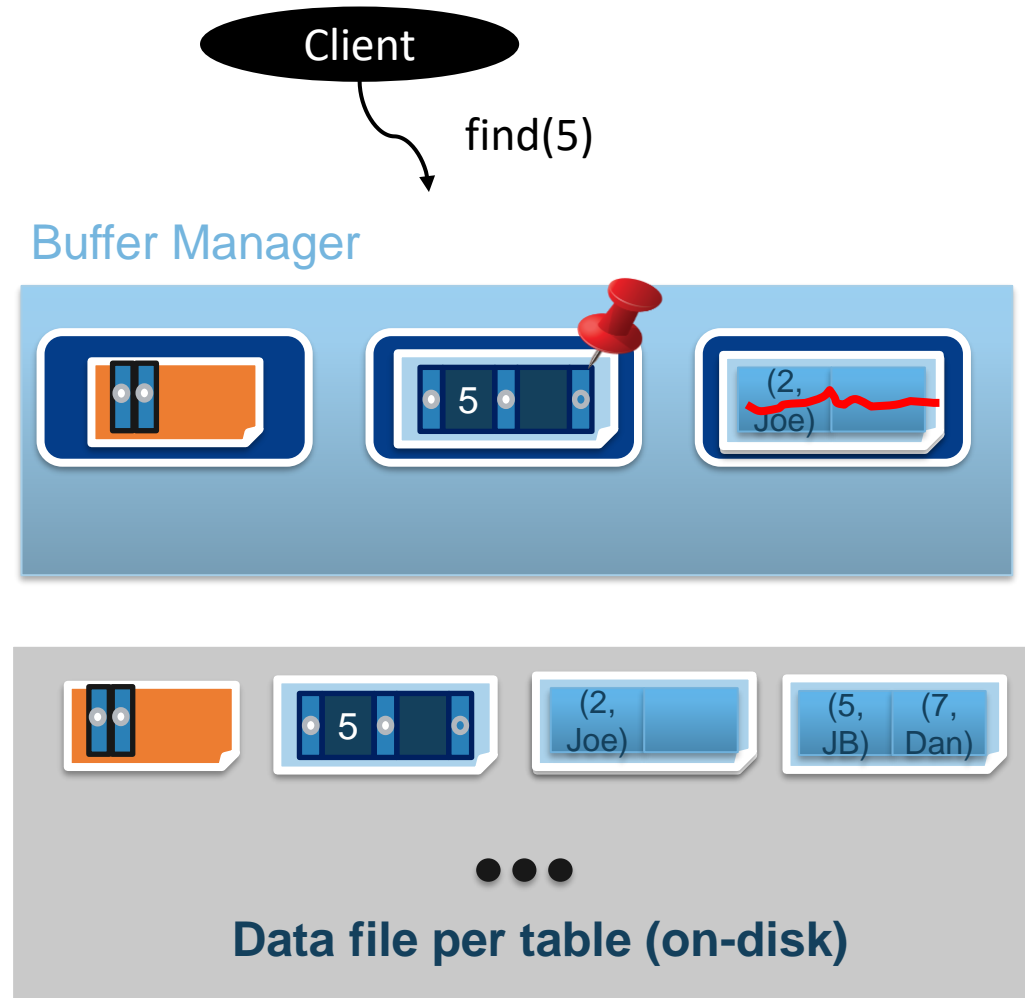
Buffer Management

- A modification operation must be handled within the in-memory buffer. You must read the page to buffer before performing any modifications.
 - Do NOT perform direct reads/writes from/to disk.
- Example (figure on the right)
 - Performing "delete key 3" will modify the page in buffer and mark the corresponding page dirty (as crossed-out in red). Notice that the on-disk page is not accessed directly.



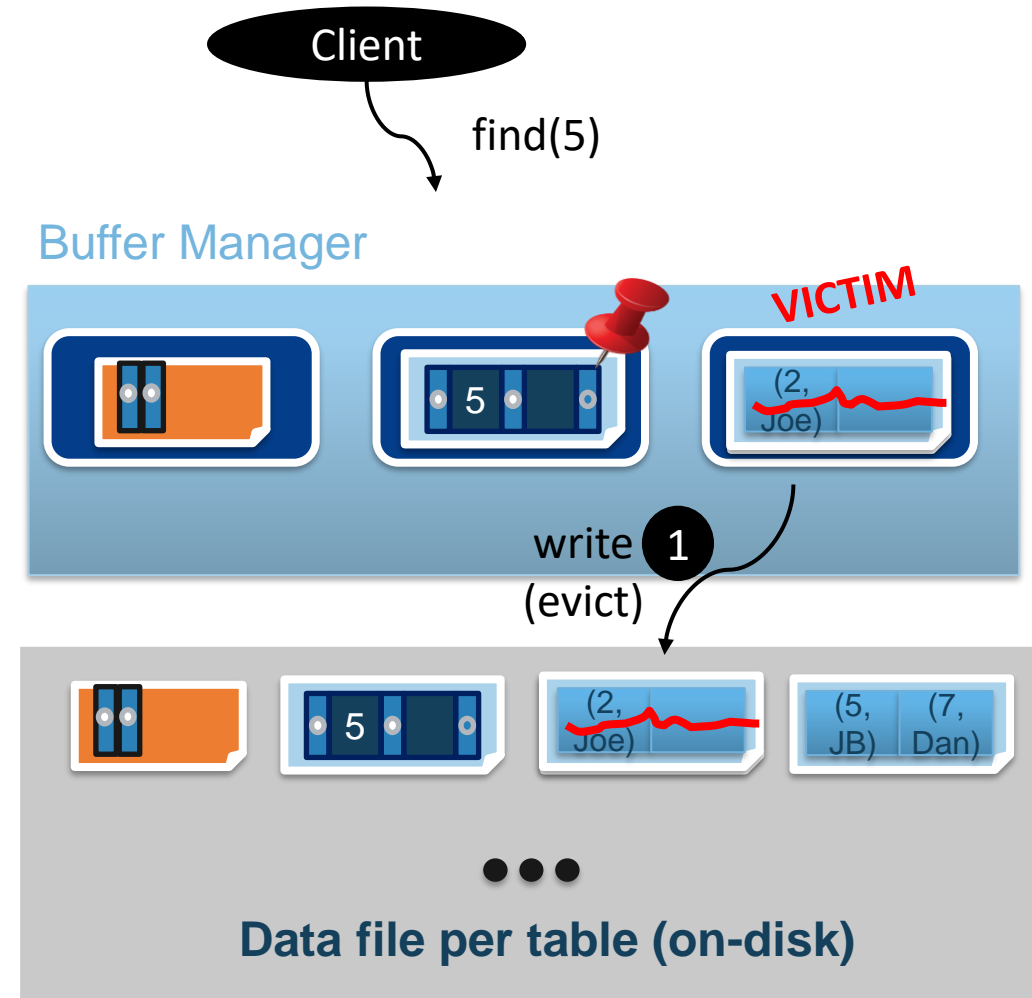
Buffer Management

- A dirty page is written to disk when the page is selected as a victim page for eviction by the LRU policy.
- Example (continue...)
 - "find key 5" attempts to read the root page.
 - The root page is found in the buffer (cache-hit)



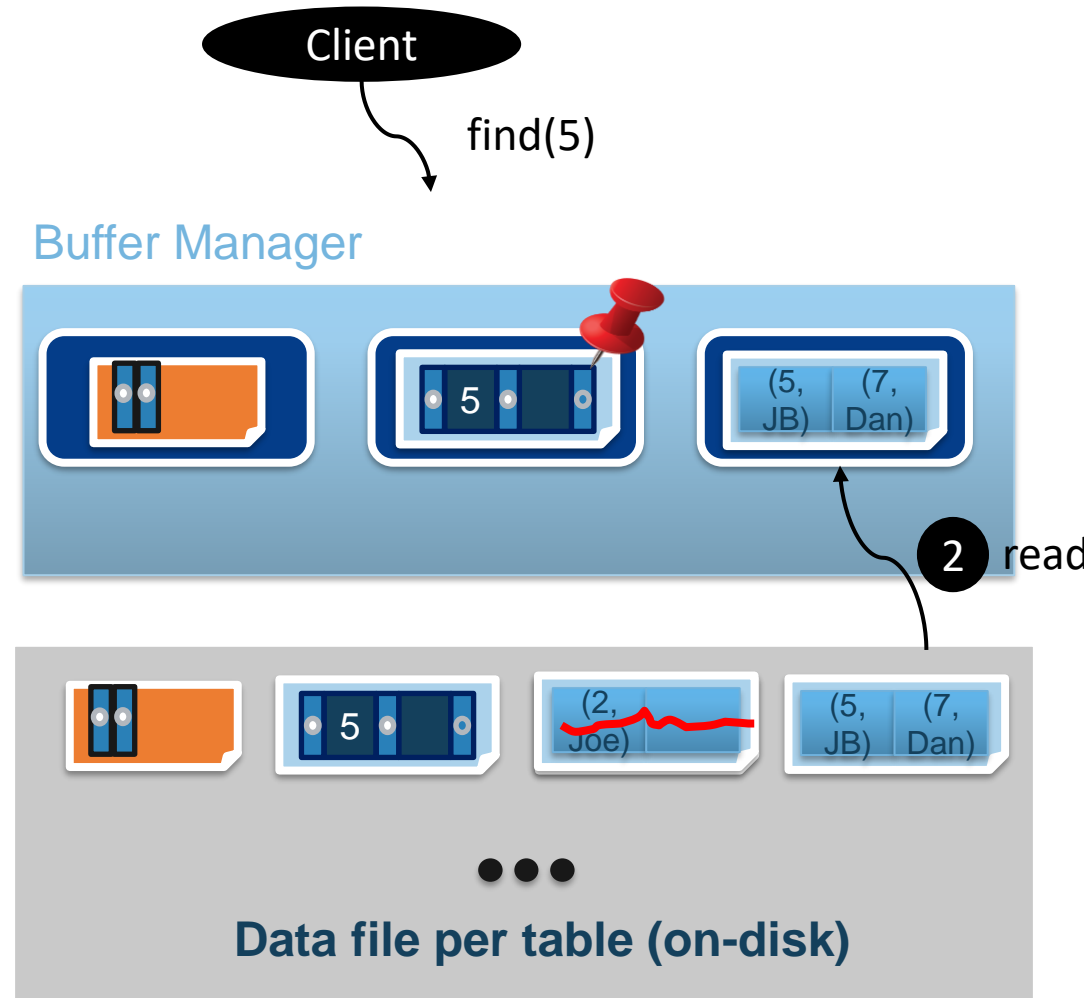
Buffer Management

- Example (continue...)
 - Traversing to the leaf, attempts to read the leaf page (with keys 5, 7). However the page is not within the buffer, cache-miss occurs.
 - The buffer manager tries to fetch the page from disk but the buffer is full.
 - The root page is evicted since it's (A) marked as dirty from the previous delete and (B) selected by the LRU policy as victim page. ①



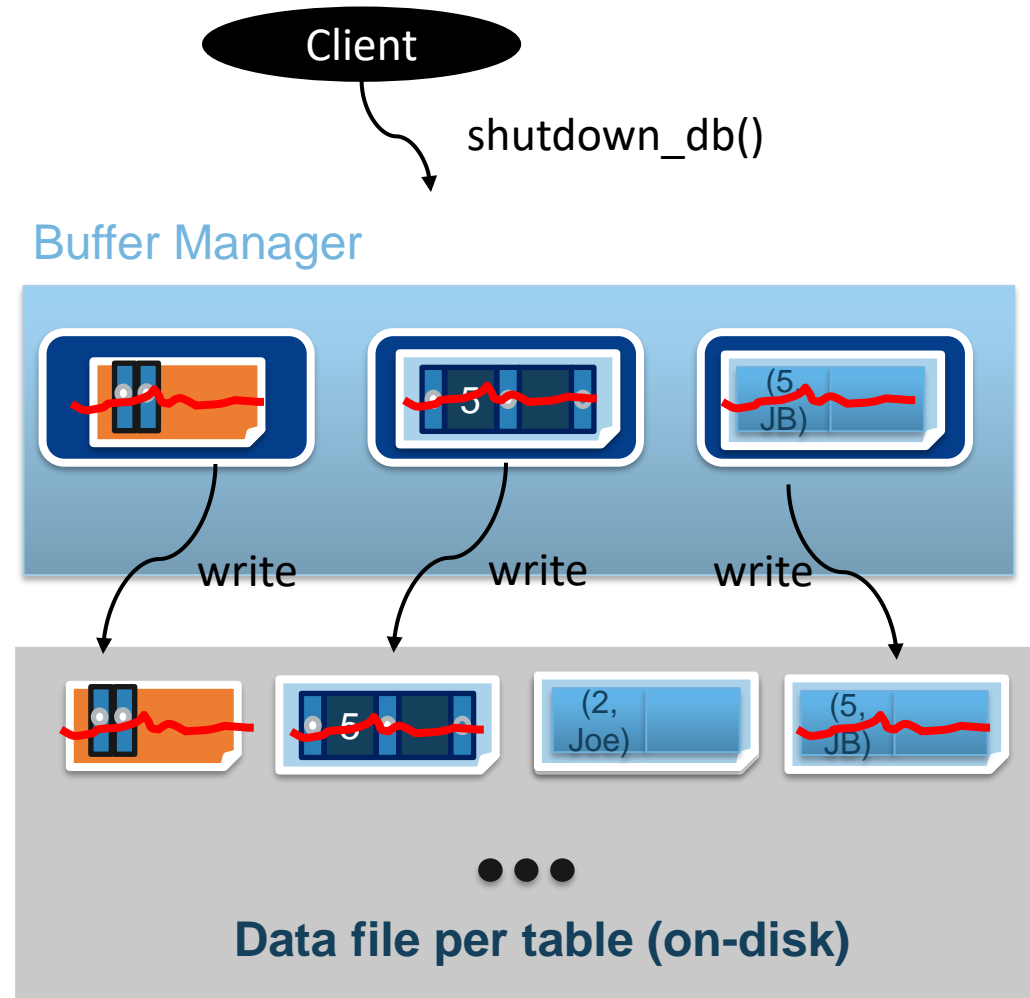
Buffer Management

- Example (continue...)
 - Traversing to the leaf, attempts to read the leaf page (with keys 5, 7). However the page is not within the buffer, cache-miss occurs.
 - The buffer manager tries to fetch the page from disk but the buffer is full.
 - The root page is evicted since it's (A) marked as dirty from the previous delete and (B) selected by the LRU policy as victim page. ①
 - Now, the corresponding leaf page can be read from disk to buffer. ②



Buffer Management

- shutdown_db() writes out all dirty buffer blocks to disk.
- This command can provide synchronous semantic (durability) to the user, but lose performance.



Submission

- Implement an in-memory buffer manager and submit a report about your design and implementation on Wiki.
 - Deadline: Nov 1 11:59pm
- We will score your project based on the last commit before deadline. Submissions afterwards will not be accepted.

Thank you
